

Chance Constrained Optimization with Kernel Density Estimation

Master-Thesis

for the degree of

Master of Science (M.Sc.)

Data Science

in the Technischen Fakultät at
Friedrich-Alexander-Universität Erlangen-Nürnberg

vorgelegt am Datum

from **Syed Ansab Waqar Gillani**

Betreuer/in: Prof. Dr. Michael Schuster



Friedrich-Alexander-Universität
Naturwissenschaftliche Fakultät

Contents

1	The Problem of Optimization Under Uncertainty	1
1.1	Where Deterministic Optimization Falls Short	1
1.1.1	The mean-value approach and the flaw of averages	1
1.1.2	Four domains where uncertainty bites	2
1.1.3	A worked example: the cost of trusting the mean	3
1.1.4	A near-universal failure rate	4
1.2	Two Extremes and the Gap Between Them	5
1.2.1	Ignoring uncertainty and worst-case robustness	5
1.2.2	When each extreme is the right tool	6
1.2.3	Chance constraints as the unifying middle ground	7
1.3	What Chance Constraints Are and Why They Are Hard	8
1.3.1	Individual and joint chance constraints	8
1.3.2	Three structural sources of difficulty	9
1.4	The Case for Nonparametric Estimation	12
1.4.1	The parametric assumption and its routine failure	12
1.4.2	What a wrong distribution does to the guarantee	13
1.4.3	Kernel density estimation and its fit to optimization	14
1.4.4	The price of going nonparametric	15
1.5	The Residual-Space Insight	15
1.5.1	From uncertainty space to residual space	15
1.5.2	Three dividends of the change of variables	16
1.5.3	Joint constraints, bandwidth, and the safety guarantee	17
1.6	Thesis Scope and Positioning	18
1.6.1	Disciplinary positioning and anchor literature	18
1.6.2	Research questions and roadmap	19
	Literatur	24

Lebenslauf

26

1 The Problem of Optimization Under Uncertainty

Every optimization model is a promise. It promises that if the decision it recommends is implemented, the constraints of the system will hold and the objective will be as good as the mathematics claims. That promise is honest only if the numbers inside the model are honest, and in almost every real deployment they are not. The demand a grid operator schedules against is a forecast, not a fact. The thrust a lander commands is corrupted by disturbances no sensor fully resolves. The returns a portfolio is built on are realizations of a process nobody has ever observed in closed form. This chapter establishes precisely what goes wrong when uncertainty is ignored, why the two classical remedies, ignoring it and armoring against the worst case, both fail in practice, and why the chance constraint, a constraint required to hold with a prescribed probability, is the correct middle ground. It then exposes the three structural reasons chance constraints are computationally hard, builds the case for estimating the relevant probability distributions from data rather than assuming them, and introduces the single conceptual shift on which the entire thesis rests: moving the estimation problem from the space of the uncertainty to the space of the constraint residual. The chapter closes by fixing the scope, the anchor literature, and the research questions of the thesis.

1.1 Where Deterministic Optimization Falls Short

1.1.1 The mean-value approach and the flaw of averages

The standard nonlinear program

$$\underset{x \in X}{\text{minimize}} \ f(x) \quad \text{subject to} \quad g(x, \xi) \leq 0 \tag{1}$$

treats the parameter vector ξ as a known constant. Decades of solver development have made problems of the form Eq. (1) routinely solvable at scale, and this success has had a quiet side effect: the deterministic template is applied by default, even when ξ is plainly a random quantity. The usual workaround is to replace ξ with a point estimate, most often its mean or a nominal forecast, and solve the resulting deterministic problem. This is the mean-value approach, and it fails in a way that is systematic rather than occasional.

The failure mechanism is worth stating carefully, because it is the engine behind every

example in this subsection. An optimizer pushes its solution to the boundary of the feasible region; that is what optimality means whenever a constraint is binding. If the constraint $g(x, \xi) \leq 0$ is active at the mean-value optimum x_{mv}^* , then the residual $g(x_{mv}^*, \xi)$ sits exactly at zero when ξ equals its mean. The moment ξ is allowed to fluctuate around that mean, roughly half of its realizations push the residual positive and the constraint is violated. The optimizer does not merely tolerate risk; it actively seeks out the configuration in which risk is maximal, because that configuration is where the objective is best. Optimization and uncertainty are adversaries by construction, a phenomenon known in the planning literature as the *flaw of averages* (Savage, 2009): a plan built on average inputs is not an average plan, it is a fragile one.

1.1.2 Four domains where uncertainty bites

Four application domains make this concrete. They are chosen deliberately, because each reappears later in the thesis as a benchmark or a case study, and because together they span the distributional pathologies that motivate the nonparametric approach of Section 1.4.

Autonomous lunar landing. A powered-descent guidance problem asks for a thrust profile that brings a lander from orbit to a designated site while respecting bounds on velocity, glide slope, and fuel. The dynamics are integrated forward under thrust perturbations and imperfectly known gravitational and terrain parameters. Keil u. a. (2021) show that disturbance distributions in this setting are naturally non-Gaussian; a mixture arises, for example, when a thruster operates in one of two regimes. A trajectory optimized for nominal thrust grazes the velocity constraint at touchdown, so under perturbation the lander exceeds its safe impact speed on a large fraction of descents. There is no second attempt at a lunar landing. Chance-constrained formulations of exactly this problem are the application anchor of the thesis (Keil u. a., 2020, 2021; Blackmore u. a., 2010).

Energy dispatch under renewable generation. A system operator commits generation a day ahead against uncertain demand and uncertain wind and solar output. Forecast errors for renewable infeed are well documented to be skewed and heavy-tailed rather than Gaussian (Hodge u. Milligan, 2011; Bienstock u. a., 2014). A dispatch that balances the forecast exactly is short of energy whenever net demand lands above its forecast, which by construction happens with probability near one half, and reserve

activations, load shedding, or price spikes follow. This domain has produced some of the strongest data-driven chance-constraint literature (Ciftci u. a., 2019; Wu u. a., 2021; Wu u. Kargarian, 2023).

Chemical process control. A reactor recipe is optimized for yield subject to temperature and purity limits, but the composition of the feed material varies from batch to batch. Running the process at the deterministic optimum places the temperature constraint on the boundary, and ordinary feedstock variation then produces off-specification batches at a steady rate. Process systems engineering was, in fact, the field where data-driven kernel-smoothed chance constraints were first formulated (Calfa u. a., 2015).

Portfolio allocation. An investor maximizes expected return subject to a cap on the probability of a large loss. Asset returns exhibit heavy tails and volatility clustering; these are among the most robust stylized facts in empirical finance (Cont, 2001). A portfolio optimized as if returns were fixed at their historical means concentrates in the highest-mean assets and breaches its loss cap far more often than any Gaussian calculation predicts (Rockafellar u. Uryasev, 2000; Liu u. a., 2022).

1.1.3 A worked example: the cost of trusting the mean

The mechanism is easiest to see in a problem small enough to solve by inspection. Consider a single-bus dispatch model. A controllable resource of capacity x megawatts must cover an uncertain net demand ξ , and capacity is costly, so the planner solves

$$\underset{x \geq 0}{\text{minimize}} \ c x \quad \text{subject to} \quad \xi - x \leq 0. \quad (2)$$

Let net demand follow a moderately skewed lognormal law with mean $\bar{\xi} \approx 102$ MW, a realistic shape for net load once renewable infeed is subtracted (Hodge u. Milligan, 2011). The mean-value method replaces ξ by $\bar{\xi}$ and returns $x_{\text{mv}}^* = \bar{\xi}$: build exactly average capacity. Figure 1(a) shows what this decision looks like against the actual demand density. The constraint is violated on every day that demand exceeds its mean, and for this distribution that is 46% of all days. A plan that appears exact on paper fails almost every other day in operation.

The repair is equally easy to read off the figure. If the planner is willing to accept a shortfall on at most 5% of days, the correct decision is the 95% quantile of demand, $x = q_{0.95} \approx 134$ MW, about 32% more capacity than the mean-value plan. Figure 1(b)




Figure 1: The mean-value fallacy in a one-dimensional dispatch problem with lognormal net demand (Monte Carlo with 2×10^5 samples). (a) The mean-value plan $x = \bar{\xi} \approx 102$ MW leaves the shaded shortfall region uncovered, so the constraint $\xi \leq x$ fails on 46% of days; covering demand with 95% reliability requires $x = q_{0.95} \approx 134$ MW. (b) Required capacity as a function of the accepted shortfall probability ε . Reliability is bought at an accelerating price as $\varepsilon \rightarrow 0$; the chance constraint makes this price explicit and lets the planner choose a point on the curve deliberately.

generalizes this observation into the central trade-off of the thesis: required capacity, and hence cost, rises steeply as the accepted violation probability ε shrinks, and rises without bound as $\varepsilon \rightarrow 0$ because the lognormal has unbounded support. Reliability is a commodity with a price, and the role of a properly formulated stochastic optimization problem is to make that price visible and chooseable rather than implicit and accidental.

1.1.4 A near-universal failure rate

One might hope that the 46% value is an artifact of the simple model. It is not. Figure 2 repeats the experiment for stylized versions of three of the domains above: dispatch under skewed demand, a landing-corridor constraint under bimodal thrust disturbance, and a portfolio loss cap under heavy-tailed (Student- t) returns. In every case the mean-value solution violates its constraint between 46% and 50% of the time, an order of magnitude beyond any tolerable design target. The violation frequency hovers near one half precisely because of the boundary mechanism described above: wherever the constraint is binding at the deterministic optimum, the median of the residual sits at or near zero, regardless of the application. The number is not a property of energy systems or of finance; it is a property of optimization itself.




Figure 2: Empirical violation frequency of the mean-value solution in three stylized problems drawn from the application domains of this thesis (Monte Carlo estimates). A typical engineering target of $\varepsilon = 5\%$ is shown for scale. The near-universal value close to 50% reflects the boundary mechanism: optimizers drive binding constraints to the edge of feasibility, where half of all perturbations cause failure.

1.2 Two Extremes and the Gap Between Them

1.2.1 Ignoring uncertainty and worst-case robustness

Once the inadequacy of the mean-value approach is accepted, the conversation moves to what should replace it. Two answers dominate, and they sit at opposite ends of a spectrum.

The first answer is to ignore uncertainty entirely, which is the mean-value approach already dismantled in Section 1.1. Its defining property deserves restating in the language used for the rest of the thesis: the deterministic solution carries no reliability statement of any kind. Nothing in the formulation Eq. (1) bounds the probability of failure; Fig. 2 shows that this probability is typically near one half. Whatever reliability a deployed mean-value solution exhibits is an accident of safety margins added by hand, outside the optimization, by engineers who do not trust the model. Hand-tuned margins are a tacit admission that the optimization solved the wrong problem.

The second answer is worst-case robust optimization (Soyster, 1973; Ben-Tal u. a., 2009). The uncertain parameter is confined to an uncertainty set Ξ , and the constraint is enforced for every element of that set:

$$g(x, \xi) \leq 0 \quad \text{for all } \xi \in \Xi. \quad (3)$$




Figure 3: The reliability spectrum for the dispatch example of Fig. 1. The deterministic mean-value plan is cheap and offers no guarantee; the worst-case robust plan is guaranteed and expensive, and is not even well defined under unbounded support. Chance constraints parameterize the entire curve between the extremes: the decision maker selects a reliability level $1 - \varepsilon$, and the optimization returns the cheapest decision that honors it.

Robust optimization is a mature and elegant theory; for many set geometries, Eq. (3) admits exact tractable reformulations, and when failure is truly never acceptable it is the correct model. Its weakness is the obverse of its strength: protection is purchased against every point of Ξ , including points that are vanishingly unlikely or physically implausible, and the objective pays for all of them. In the dispatch example, the demand distribution has unbounded support, so a literal worst-case formulation has no feasible solution at all; the set Ξ must be truncated by fiat, and the answer then depends sensitively on a truncation the data does not specify. Even with a defensible truncation, the protection level is an all-or-nothing property of the set, not a dial the decision maker controls. Empirically, robust solutions in power systems and aerospace routinely sacrifice a large fraction of attainable objective value to guard against scenarios whose combined probability is far below any operational concern (Ben-Tal u. a., 2009; Keil u. a., 2021).

1.2.2 When each extreme is the right tool

Fairness demands an acknowledgment before the middle ground is claimed: each extreme has a legitimate home. The mean-value approach is correct when uncertainty is genuinely negligible relative to the constraint margins, or when violations carry only a linear cost that the objective already prices in, as in inventory models where a shortfall is simply an expense. Worst-case robustness is correct when the uncertainty




Figure 4: Three responses to constraint uncertainty, shown on the density of the residual $Y_x = g(x, \xi)$; the constraint is satisfied for $Y_x \leq 0$. **(a)** The mean-value solution satisfies the constraint at $\bar{\xi}$ but leaves an unquantified violation mass (red). **(b)** The worst-case robust solution eliminates all violations by pushing the distribution well clear of zero, at the cost of large unused margin. **(c)** The chance-constrained solution places exactly the permitted fraction ε of probability mass beyond the boundary, making the accepted risk explicit and controllable.

set is small, well characterized, and physically bounded, and when a single failure is catastrophic and irreversible: structural safety factors live here for good reason. The problems this thesis targets sit in neither home. Their uncertainties are large, data-described, and unbounded or effectively so, and their failures are costly but tolerable at a small, explicit rate. A lander program accepts a quantified risk per descent; a grid code tolerates a bounded loss-of-load expectation per year. For this regime, neither extreme expresses what the decision maker actually wants.

1.2.3 Chance constraints as the unifying middle ground

Between a method with no guarantee and a method that guarantees too much lies the formulation this thesis is about. A chance constraint, introduced by Charnes u. Cooper (1959) and developed into a full branch of stochastic programming by Prékopa (1995), requires the constraint to hold with a prescribed probability:

$$\mathbb{P}(g(x, \xi) \leq 0) \geq 1 - \varepsilon. \quad (4)$$

The risk tolerance $\varepsilon \in (0, 1)$ is small, typically between 0.001 and 0.1 depending on the stakes, and $1 - \varepsilon$ is the reliability level. Figure 3 places the three formulations on a single axis for the dispatch example. The chance constraint is not a compromise between the extremes; it is a strict generalization of both. As $\varepsilon \rightarrow 0$ it recovers worst-case protection over the support of ξ , and at $\varepsilon = 0.5$ it roughly reproduces the mean-

value behavior for symmetric residuals. Every intermediate point is available, and each carries an explicit, auditable meaning: the system fails at most a fraction ε of the time. This is exactly the language in which engineering requirements are written. Grid codes prescribe loss-of-load expectations, flight rules prescribe acceptable failure probabilities per descent, and risk mandates prescribe loss probabilities per quarter. The chance constraint is the optimization-native translation of a reliability requirement, which is why it is the right framing for most engineering and operational problems, and why the difficulty of solving it, the subject of Section 1.3, has remained an active research frontier for over six decades.

Figure 4 makes the three positions concrete. Each panel shows the same bimodal residual density $Y_x = g(x, \xi)$, with the constraint boundary at zero; the two modes reflect the non-Gaussian disturbance shapes that motivate the nonparametric approach developed from subsection 1.4 onward

1.3 What Chance Constraints Are and Why They Are Hard

1.3.1 Individual and joint chance constraints

This subsection fixes notation used throughout the thesis and isolates the three structural sources of computational difficulty. Let $x \in X \subseteq \mathbb{R}^n$ be the decision vector, let ξ be a random vector on a probability space $(\Omega, \mathcal{F}, \mathbb{P})$ taking values in $\mathbb{R}^d$, and let $g : \mathbb{R}^n \times \mathbb{R}^d \rightarrow \mathbb{R}$ be a constraint function. The *individual chance constraint* is

$$\mathbb{P}(g(x, \xi) \leq 0) \geq 1 - \varepsilon. \quad (5)$$

When the system carries several uncertain requirements $g_1, \dots, g_R$ that must hold simultaneously, the correct model is the *joint chance constraint*

$$\mathbb{P}(g_r(x, \xi) \leq 0, r = 1, \dots, R) \geq 1 - \varepsilon, \quad (6)$$

which bounds the probability that *all* requirements hold at once. The distinction matters: enforcing each g_r individually at level $1 - \varepsilon$ does not bound the joint failure probability, while the classical repair via Boole's inequality, splitting the budget as $\sum_r \varepsilon_r \leq \varepsilon$, is sound but conservative (Prékopa, 1995; Keil u. a., 2021). The full chance-constrained program studied in this thesis is

$$\underset{x \in X}{\text{minimize}} \ f(x) \quad \text{subject to} \quad \mathbb{P}(g(x, \xi) \leq 0) \geq 1 - \varepsilon. \quad (7)$$

A word on the role of ε is warranted, because it is the one quantity in Eq. (7) that mathematics cannot supply. The risk tolerance is a design parameter set by the consequences of failure and by external requirements: a 5% daily shortfall tolerance may be acceptable for a microgrid with backup supply (Ciftci u. a., 2019), while a descent constraint may demand $\varepsilon \leq 10^{-2}$ or far less (Blackmore u. a., 2010). Two facts about this parameter shape everything downstream. First, the difficulty of the problem grows as ε shrinks, because the constraint then probes ever deeper into the tail of a distribution that must be known or estimated ever more accurately there. Second, ε is a budget, and for joint constraints Eq. (6) the budget must cover all R requirements simultaneously, which is what makes the joint form strictly harder than the individual one. Throughout the thesis, ε is treated as given and the question is how to honor it efficiently and honestly.

1.3.2 Three structural sources of difficulty

Writing Eq. (7) is easy; solving it is not, and the obstruction is threefold. Each difficulty below is independent of the others, and any viable solution method must answer all three. The pattern repeats through the entire methodological literature, so these three points are the lens through which every method in this thesis, including the KDE reformulation itself, should be judged.

Difficulty 1: the feasible set is generally non-convex. Convexity of g in x does not transfer to the probabilistic feasible set $\{x : \mathbb{P}(g(x, \xi) \leq 0) \geq 1 - \varepsilon\}$. A minimal example makes the geometry vivid. Take a linear constraint $\xi^\top x \leq 1$ where ξ takes one of two values $\xi^{(1)}, \xi^{(2)}$ with probability one half each, and set $\varepsilon = 0.5$. The chance constraint then requires $\xi^\top x \leq 1$ for *at least one* of the two scenarios, so the feasible set is the union of two half-spaces. Figure 5(a) shows the result: two points, each satisfying a different scenario, are individually feasible, while their midpoint satisfies neither and is infeasible. A union of convex sets is not convex, and this toy construction scales up: with general discrete or continuous distributions the feasible set can shatter into combinatorially many pieces. Convexity survives only under special structure, most famously when the distribution of ξ is log-concave and g has suitable form, the celebrated theory of Prékopa (1995), or when ξ is Gaussian and the constraint is linear in ξ , in which case Eq. (5) collapses to a second-order cone constraint (Charnes u. Cooper, 1963; Shapiro u. a., 2014). Outside these islands, non-convexity is the rule, and global optimality claims must be abandoned in favor of high-quality local solutions, which is precisely the regime in which nonlinear programming solvers operate.




Figure 5: The two geometric faces of chance-constraint hardness. (a) For a two-point distribution of ξ and $\varepsilon = 0.5$, the feasible set of $\mathbb{P}(\xi^\top x \leq 1) \geq 0.5$ is the union of two half-spaces (shaded). The marked endpoints are feasible while their midpoint is not, certifying non-convexity. (b) The contribution of a single sample to the probability is the discontinuous indicator $\mathbf{1}\{y \leq 0\}$ of the residual; its derivative is zero almost everywhere and undefined at the jump. Smooth surrogates obtained by integrating a kernel, shown for two bandwidths h , restore the differentiability that gradient-based solvers require, at the price of a bias controlled by h . This single picture contains the central tension of the thesis.

Difficulty 2: the probability has no closed form. Define the parameterized failure region $M(x) = \{\xi : g(x, \xi) > 0\}$. Evaluating the constraint in Eq. (7) means computing

$$\mathbb{P}(g(x, \xi) \leq 0) = 1 - \int_{M(x)} \rho_\xi(\xi) d\xi, \quad (8)$$

an integral of the density ρ_ξ over a region whose shape changes with every candidate decision x . For nonlinear g the region $M(x)$ is curved and possibly disconnected, the integral is multidimensional, and a closed form exists only in textbook cases. Worse, in the data-driven settings that motivate this thesis the density ρ_ξ itself is unknown; all that exists is a finite collection of samples from simulation, historical records, or experiments. The constraint function that the optimizer must evaluate, thousands of times, at thousands of points x , is therefore an unknown integral of an unknown density over an intractable region. Any practical method replaces Eq. (8) with an approximation, and the quality of that approximation governs everything downstream.

Difficulty 3: the probability is built from a discontinuous function. Rewriting the probability as an expectation,

$$\mathbb{P}(g(x, \xi) \leq 0) = \mathbb{E}[\mathbf{1}\{g(x, \xi) \leq 0\}], \quad (9)$$

exposes the third obstruction: the indicator $\mathbf{1}\{\cdot\}$ jumps from one to zero at the constraint boundary, as plotted in Fig. 5(b). The natural sample-based estimate of Eq. (9), count the fraction of samples satisfying the constraint, inherits this discontinuity. As a function of x it is piecewise constant: its gradient is zero almost everywhere and undefined on a measure-zero set, which is the worst possible input for the sequential quadratic programming and interior-point methods that dominate nonlinear optimization (Shapiro u. a., 2014). The established escape routes each pay a recognizable price. Mixed-integer reformulations of the counting constraint are exact but combinatorial and scale poorly (Luedtke u. Ahmed, 2008). The scenario approach enforces the constraint on all N sampled scenarios, regaining convexity and even finite-sample feasibility certificates for convex problems (Calafiore u. Campi, 2006; Campi u. Garatti, 2008), but it cannot exploit the slack that ε permits and grows conservative as N increases. Conservative convex surrogates such as conditional value-at-risk restrict the feasible set from inside (Rockafellar u. Uryasev, 2000; Nemirovski u. Shapiro, 2006), guaranteeing feasibility at a cost in objective value. Smoothing methods, the family to which this thesis belongs, replace the indicator with a differentiable approximation (Peña-Ordieres u. a., 2020), restoring the gradients that nonlinear solvers need while introducing a controlled bias, visible in Fig. 5(b) as the gap between the step and the smooth curves.

The combination of the three difficulties, and not any one of them alone, is what makes Eq. (7) fundamentally harder than a deterministic NLP. Non-convexity rules out global guarantees, the missing closed form rules out exact evaluation, and non-smoothness rules out naive sample averaging inside derivative-based solvers. Every approximation method discussed in this thesis is an attempt to buy back tractability by conceding something: exactness, generality, or a margin of conservatism. The contribution of the KDE-based reformulation is a specific and unusually favorable bargain on all three fronts, and the remainder of this chapter explains why.




Figure 6: The three structural sources of difficulty in chance-constrained programming: (a) the feasible set is non-convex, (b) the constraint probability has no closed form because the safe region moves with x , and (c) the probability function is flat almost everywhere and non-differentiable at the boundary.

1.4 The Case for Nonparametric Estimation

1.4.1 The parametric assumption and its routine failure

Among the classical escape routes from Section 1.3, the oldest and still most widely used is parametric: assume the law of ξ belongs to a named family, almost always the Gaussian, and exploit the structure that follows. Under Gaussian ξ and a constraint linear in ξ , the chance constraint Eq. (5) reduces exactly to a deterministic second-order cone constraint involving the mean, the covariance, and a quantile of the standard normal (Charnes u. Cooper, 1963; Shapiro u. a., 2014). The reformulation is convex, smooth, and supported by every modern conic solver. When the Gaussian assumption is true, nothing in this thesis improves on it, and the experimental chapters use it as the reference baseline for exactly that reason.

The assumption, however, is routinely false in the very domains where chance constraints matter most, and the evidence is not anecdotal. Wind-power forecast errors are skewed and heavy-tailed, with documented deviations from normality that grow in the tails, which is precisely where a chance constraint lives (Hodge u. Milligan, 2011; Wan u. a., 2020). Equity and commodity returns exhibit heavy tails and aggregational Gaussianity only at long horizons; at decision-relevant horizons their kurtosis is far above three (Cont, 2001). Thrust disturbances and navigation errors in powered descent arise from mixtures of operating regimes and are naturally multimodal (Keil u. a., 2021). Process feedstocks vary by supplier and batch in ways no single unimodal family captures (Calfa u. a., 2015). The Gaussian is chosen for the convenience of the




Figure 7: What a wrong distributional assumption does to a chance constraint. (a) A bimodal disturbance law, of the kind produced by two-regime actuator behavior, against its best Gaussian fit: the fit places probability mass between the modes where the true law has almost none, and misplaces both shoulders. (b) A skewed forecast-error law (centered Gamma) on a logarithmic density scale. Beyond the marked threshold the Gaussian fit reports 0.23% tail mass where the true law carries 1.7%, an underestimate by a factor of about seven. A planner who certifies a 0.5% risk level with the Gaussian model deploys a system whose true failure rate is more than three times the certificate.

analyst, not the behavior of the system.

1.4.2 What a wrong distribution does to the guarantee

The consequences of the mismatch are asymmetric and quantifiable. Figure 7 dissects both failure modes. Panel (a) shows a bimodal disturbance law next to its moment-matched Gaussian fit: the fit invents probability where there is none and erases the two modes that actually govern system behavior. Panel (b) shows the more dangerous case on a logarithmic scale. For a skewed forecast-error law, the Gaussian fit underestimates the mass beyond a critical threshold by a factor of seven: it reports 0.23% where the truth is 1.7%. A decision certified to a 0.5% failure probability under the Gaussian model in fact fails more than three times as often as certified. The error can also point the other way: when the true law is lighter-tailed than the fit, the parametric constraint is needlessly conservative and pays an invisible objective penalty. Falsely optimistic or falsely conservative, the direction is decided by distributional geometry the analyst assumed away, and the certificate that the chance constraint was supposed to provide, $\mathbb{P}(\text{failure}) \leq \varepsilon$, silently becomes false. A reliability guarantee that depends on an unverified distributional assumption is not a guarantee; it is a hope with a confidence interval of unknown width.

1.4.3 Kernel density estimation and its fit to optimization

The alternative pursued in this thesis is to estimate the distribution from the data instead of assuming it. Kernel density estimation, the canonical nonparametric estimator introduced by Rosenblatt (1956) and Parzen (1962), reconstructs a density from samples $y_1, \dots, y_N$ as

$$\hat{\rho}_h(y) = \frac{1}{Nh} \sum_{j=1}^N K\left(\frac{y - y_j}{h}\right), \quad (10)$$

where the kernel K is a smooth density and the bandwidth $h > 0$ sets the degree of smoothing. The estimator makes no shape commitment: given enough data it recovers skewness, multimodality, and heavy tails alike, with well-understood consistency theory and practical bandwidth rules (Silverman, 1986; Devroye u. Györfi, 1985). Three properties make KDE specifically suited to the optimization context, rather than merely to statistics. First, $\hat{\rho}_h$ is as smooth as its kernel, so quantities built from it are differentiable, which addresses Difficulty 3 head-on. Second, its integral is available in closed form whenever the kernel’s cumulative distribution function is, which addresses the evaluation half of Difficulty 2. Third, it is distribution-free, which removes the silent assumption dismantled above.

This is not a speculative combination. A coherent body of work already demonstrates KDE-based chance constraints across the application domains of Section 1.1: kernel-smoothed individual and joint chance constraints in process planning (Calfa u. a., 2015), KDE-reformulated optimal control for launcher ascent and the Goddard problem (Caillaud u. a., 2018), biased-KDE optimal control for soft lunar landing (Keil u. a., 2020, 2021; Keil u. Rao, 2021), adaptive-KDE microgrid energy management (Ciftci, 2017; Ciftci u. a., 2019; Babić u. a., 2023), multivariate-KDE joint chance constraints for economic dispatch and power-system scheduling (Wu u. a., 2021; Wu u. Kargarian, 2023), KDE-driven distributionally robust load shedding (Qi u. a., 2026), and KDE-based tail-risk portfolio optimization (Liu u. a., 2022). Adjacent nonparametric formulations reinforce the same direction in robotics and forecasting (Gopalakrishnan u. a., 2018; Wan u. a., 2020; Wang u. a., 2021). What this literature has lacked, until recently, is a unifying account connecting the engineering constructions to convergence theory; supplying and stress-testing that connection is the work of this thesis.

1.4.4 The price of going nonparametric

A nonparametric estimator does not come free, and the costs are named now because they structure the experimental program of Chapters 7. KDE suffers the curse of dimensionality: its convergence rate degrades rapidly as the dimension of the estimated variable grows (Silverman, 1986). The bandwidth h is a genuine hyperparameter whose role inside an optimization problem, where it shapes the feasible region and the gradients the solver sees, is more delicate than its role in density estimation. And a plain KDE smooths in both directions, so it can overestimate the safe probability and thereby violate the very bound it is meant to enforce. Each of these costs has a specific answer in this thesis: the first is neutralized by the residual-space construction of Section 1.5, the second is treated as a first-class experimental question, and the third is resolved by the biased-KDE construction of Keil u. a. (2021).

1.5 The Residual-Space Insight

1.5.1 From uncertainty space to residual space

Everything assembled so far converges on a single conceptual move, simple to state and consequential enough to carry the thesis. The naive way to use KDE in Eq. (7) is to estimate the density of the uncertainty ξ itself and then integrate the estimate over the safe region $\mathbb{R}^d \setminus M(x)$. This inherits both halves of Difficulty 2: the estimation runs in d dimensions, exactly where the curse of dimensionality bites, and the integration region still deforms with every candidate x . The correct object to estimate is not the uncertainty. It is the *constraint residual*

$$Y(x) = g(x, \xi), \tag{11}$$

the scalar random variable obtained by pushing the uncertainty through the constraint function at the current decision. The residual is the quantity the constraint actually speaks about: $Y(x) \leq 0$ is constraint satisfaction, by definition, and the chance constraint Eq. (5) is nothing but a statement about the cumulative distribution function of $Y(x)$ at zero,

$$\mathbb{P}(Y(x) \leq 0) = F_{Y(x)}(0) \geq 1 - \varepsilon. \tag{12}$$

Schuster (2025) makes this change of variables the foundation of the convergence theory: the probability is rewritten as an integral of the density of the *transformed*




Figure 8: The residual-space construction at a fixed candidate decision x . Samples ξ_j of the uncertainty are pushed through the constraint to give residual samples $y_j = g(x, \xi_j)$ (histogram, $N = 400$, bimodal by construction); a KDE $\hat{\rho}_{Y(x)}$ smooths them, and its integral over the fixed safe region $(-\infty, 0]$ (shaded) estimates the safe probability, here $\hat{P} \approx 0.67$. The estimation problem is one-dimensional regardless of the dimension of ξ , and the resulting constraint $\hat{F}_{Y(x)}(0) \geq 1 - \varepsilon$ is a smooth, differentiable function of x .

variable $g(x, \xi)$ over the fixed set $(-\infty, 0]$, rather than an integral of the density of ξ over the moving set $\mathbb{R}^d \setminus M(x)$.

1.5.2 Three dividends of the change of variables

The dividends of this shift are worth enumerating, because each one disarms a specific difficulty from Section 1.3.

First, *dimensionality collapses*. The uncertainty ξ may live in ten or a hundred dimensions; the residual of an individual constraint is a scalar, always, and the residuals of a joint constraint form a vector of dimension R , the number of constraint components, which in the problems of this thesis is small. KDE is being asked to do the one job it does excellently, one-dimensional density estimation from a few hundred samples, and the curse of dimensionality in d is bypassed entirely. The constraint function g itself performs the dimension reduction, and it does so without information loss for the question being asked, since Eq. (12) is an identity, not an approximation.

Second, *the integration region freezes*. The safe region in residual space is $(-\infty, 0]$ for every decision x ; all dependence on x has migrated into the samples $y_j = g(x, \xi_j)$ themselves. Integrating the kernel estimate over a fixed half-line is elementary: substituting

Eq. (10) gives the estimated CDF

$$\hat{F}_{Y(x)}(0) = \frac{1}{N} \sum_{j=1}^N \mathcal{K}\left(\frac{-g(x, \xi_j)}{h}\right), \quad \mathcal{K}(t) = \int_{-\infty}^t K(s) ds, \quad (13)$$

a finite sum of evaluations of the kernel's CDF $\mathcal{K}$, available in closed form for every kernel used in this thesis.

Third, *smoothness is restored*. Comparing Eq. (13) with the sample-average form Eq. (9) reveals the substitution at the heart of the method: the discontinuous indicator $\mathbf{1}\{g(x, \xi_j) \leq 0\}$ has been replaced by the smooth sigmoid $\mathcal{K}(-g(x, \xi_j)/h)$, exactly the curves drawn in Fig. 5(b). If g is differentiable in x , then so is $\hat{F}_{Y(x)}(0)$, with gradients obtained by the chain rule through $\mathcal{K}$. The chance constraint Eq. (5) thereby becomes the deterministic, smooth, nonlinear constraint

$$\hat{F}_{Y(x)}(0) \geq 1 - \varepsilon, \quad (14)$$

which any sequential quadratic programming, interior-point, or direct-collocation solver consumes natively. Figure 8 shows the whole construction in one frame: residual samples, the kernel estimate, the frozen safe region, and the resulting probability estimate.

1.5.3 Joint constraints, bandwidth, and the safety guarantee

The construction extends to joint chance constraints with one further idea that previews the extension studied in this thesis. For Eq. (6), the natural scalar summary of the constraint system is the worst violation, $Y(x) = \max_r g_r(x, \xi)$, since all requirements hold exactly when the maximum residual is nonpositive. The maximum preserves the one-dimensionality of the estimation problem but reintroduces a kink; smoothing it with a log-sum-exp surrogate restores differentiability at the cost of a second, controllable approximation layer. This scalar violation statistic, its interaction with the KDE layer, and the alternative of low-dimensional multivariate KDE over the residual vector form the substance of research question RQ4 and are developed in Part 3.

The reformulation Eq. (14) is where the thesis begins rather than ends, because the two costs flagged in Section 1.4 are still standing and they define the methodological arc of Part 3. The bandwidth h now controls the geometry of an optimization problem, not just the fidelity of a density plot: $h \rightarrow 0$ recovers the non-smooth indicator and its

numerical pathologies, while large h oversmooths and distorts the feasible region, and no classical bandwidth rule was derived with either effect in mind. And because an ordinary KDE can overestimate $F_{Y(x)}(0)$, the constraint Eq. (14) can certify a reliability the true system does not possess. The constructive answer is the biased KDE of Keil u. a. (2020, 2021): kernels engineered so that their integral function bounds the indicator from the safe side, making Eq. (14) a guaranteed conservative approximation of the true chance constraint while remaining smooth and bounded. The theoretical answer is the convergence framework of Schuster (2025), which establishes conditions under which solutions of the KDE-approximated problem converge to solutions of the true problem as $N \rightarrow \infty$. Engineering construction and convergence theory, developed independently, meet in residual space, and their meeting is the subject of this thesis.

1.6 Thesis Scope and Positioning

1.6.1 Disciplinary positioning and anchor literature

A thesis at the intersubsection of three fields must say plainly what it is and what it is not. This thesis is not a contribution to nonparametric statistics: it proves no new properties of kernel density estimators, and it takes the consistency theory of Parzen (1962); Silverman (1986); Devroye u. Györfi (1985) as given. It is not a contribution to control theory: it establishes no stability or recursive feasibility results for closed-loop predictive control. It is a data science and applied optimization contribution. Its object of study is the KDE-based reformulation pipeline itself, from samples through residuals through Eq. (14) into a numerical solver, and its method is to implement that pipeline rigorously, analyze its behavior, compare it against the established alternatives, and validate it across problem settings chosen to stress its assumptions.

Two anchor papers bound the territory. Keil u. a. (2021), with the earlier preprint Keil u. a. (2020) and the computational follow-up Keil u. Rao (2021), supply the constructive, engineering side: biased kernels whose integral functions bound the indicator, a safety-preserving deterministic reformulation, and a demonstration on chance-constrained soft lunar landing solved by direct collocation. Schuster (2025) supplies the analytical side: convergence theorems guaranteeing that, under verifiable conditions on the kernel, the bandwidth, and the problem, solving the estimated problem is asymptotically equivalent to solving the true one. The first anchor gives the method its practical credibility, the second its mathematical license; the thesis builds

the bridge between them and load-tests it experimentally. The surrounding evidence base spans aerospace (Caillau u. a., 2018), energy systems (Ciftci u. a., 2019; Wu u. a., 2021; Wu u. Kargarian, 2023; Qi u. a., 2026), process systems (Calfa u. a., 2015), and finance (Liu u. a., 2022), and the comparative baselines are the classical ones: scenario optimization (Calafiore u. Campi, 2006; Campi u. Garatti, 2008), sample average approximation (Luedtke u. Ahmed, 2008; Pagnoncelli u. a., 2009), parametric reformulation (Charnes u. Cooper, 1963), and convex risk surrogates (Rockafellar u. Uryasev, 2000; Nemirovski u. Shapiro, 2006).

1.6.2 Research questions and roadmap

Four research questions organize everything that follows.

- **RQ1 (formulation):** how is a chance-constrained problem reformulated through KDE over constraint residuals, for individual and joint constraints, and under what conditions is the reformulation mathematically justified?
- **RQ2 (numerical behavior):** how does the KDE reformulation compare with deterministic, scenario-based, and parametric approaches in feasibility, objective value, conservatism, runtime, and sensitivity to sample size?
- **RQ3 (bandwidth and smoothing):** how do bandwidth and smoothing choices shape feasibility, conservatism, and solver performance, given that h acts on the optimization geometry and not only on estimation error?
- **RQ4 (extension):** can the framework be strengthened through one focused extension, specifically scalar violation statistics for joint chance constraints, optionally with log-sum-exp smoothing?

The argument now proceeds in three movements. Chapter 2 supplies the technical vocabulary. Chapters 3 — 6 develop the mathematics in the order a practitioner would discover it. Chapter 7 onwards validates the framework experimentally. The destination is a defensible answer to a question every one of the systems in Section 1.1 silently poses: how safe is safe enough, and what is the cheapest decision that honestly delivers it?

References

- [Babić u. a. 2023] BABIĆ, Luka u. a.: Data-driven non-parametric chance-constrained model predictive control for microgrids energy management using small data batches. In: *Frontiers in Control Engineering* 4 (2023), S. 1237759. <http://dx.doi.org/10.3389/fcteg.2023.1237759>. – DOI 10.3389/fcteg.2023.1237759
- [Ben-Tal u. a. 2009] BEN-TAL, Aharon ; EL GHAOU, Laurent ; NEMIROVSKI, Arkadi: *Robust Optimization*. Princeton, NJ : Princeton University Press, 2009 (Princeton Series in Applied Mathematics)
- [Bienstock u. a. 2014] BIENSTOCK, Daniel ; CHERTKOV, Michael ; HARNETT, Sean: Chance-constrained optimal power flow: risk-aware network control under uncertainty. In: *SIAM Review* 56 (2014), Nr. 3, S. 461–495. <http://dx.doi.org/10.1137/130910312>. – DOI 10.1137/130910312
- [Blackmore u. a. 2010] BLACKMORE, Lars ; ONO, Masahiro ; BEKTASSOV, Askar ; WILLIAMS, Brian C.: A probabilistic particle-control approximation of chance-constrained stochastic predictive control. In: *IEEE Transactions on Robotics* 26 (2010), Nr. 3, S. 502–517. <http://dx.doi.org/10.1109/TRO.2010.2044948>. – DOI 10.1109/TRO.2010.2044948
- [Caillaud u. a. 2018] CAILLAU, Jean-Baptiste ; CERF, Max ; SASSI, Achille ; TRÉLAT, Emmanuel ; ZIDANI, Hasnaa: Solving chance constrained optimal control problems in aerospace via kernel density estimation. In: *Optimal Control Applications and Methods* 39 (2018), Nr. 5, S. 1833–1858. <http://dx.doi.org/10.1002/oca.2445>. – DOI 10.1002/oca.2445
- [Calafiore u. Campi 2006] CALAFIORE, Giuseppe C. ; CAMPI, Marco C.: The scenario approach to robust control design. In: *IEEE Transactions on Automatic Control* 51 (2006), Nr. 5, S. 742–753. <http://dx.doi.org/10.1109/TAC.2006.875041>. – DOI 10.1109/TAC.2006.875041
- [Calfa u. a. 2015] CALFA, Bruno A. ; GROSSMANN, Ignacio E. ; AGARWAL, Anshul ; BURY, Scott J. ; WASSICK, John M.: Data-driven individual and joint chance-constrained optimization via kernel smoothing. In: *Computers & Chemical Engineering* 78 (2015), S. 51–69. <http://dx.doi.org/10.1016/j.compchemeng.2015.04.012>. – DOI 10.1016/j.compchemeng.2015.04.012

- [Campi u. Garatti 2008] CAMPI, Marco C. ; GARATTI, Simone: The exact feasibility of randomized solutions of uncertain convex programs. In: *SIAM Journal on Optimization* 19 (2008), Nr. 3, S. 1211–1230. <http://dx.doi.org/10.1137/07069821X>. – DOI 10.1137/07069821X
- [Charnes u. Cooper 1959] CHARNES, Abraham ; COOPER, William W.: Chance-constrained programming. In: *Management Science* 6 (1959), Nr. 1, S. 73–79. <http://dx.doi.org/10.1287/mnsc.6.1.73>. – DOI 10.1287/mnsc.6.1.73
- [Charnes u. Cooper 1963] CHARNES, Abraham ; COOPER, William W.: Deterministic equivalents for optimizing and satisficing under chance constraints. In: *Operations Research* 11 (1963), Nr. 1, S. 18–39. <http://dx.doi.org/10.1287/opre.11.1.18>. – DOI 10.1287/opre.11.1.18
- [Ciftci 2017] CİFTCI, Okan: *Microgrid Energy Management with Flexibility Constraints: A Data-Driven Solution Method*, Louisiana State University, Diplomarbeit, 2017. http://dx.doi.org/10.31390/gradschool_theses.4351. – DOI 10.31390/gradschool_theses.4351
- [Ciftci u. a. 2019] CİFTCI, Okan ; MEHRTASH, Mahdi ; KARGARIAN, Amin: Data-driven nonparametric chance-constrained optimization for microgrid energy management. In: *IEEE Transactions on Industrial Informatics* 16 (2019), Nr. 4, S. 2447–2457. <http://dx.doi.org/10.1109/TII.2019.2932078>. – DOI 10.1109/TII.2019.2932078
- [Cont 2001] CONT, Rama: Empirical properties of asset returns: stylized facts and statistical issues. In: *Quantitative Finance* 1 (2001), Nr. 2, S. 223–236. <http://dx.doi.org/10.1080/713665670>. – DOI 10.1080/713665670
- [Devroye u. Györfi 1985] DEVROYE, Luc ; GYÖRFI, László: *Nonparametric Density Estimation: The L_1 View*. New York : Wiley, 1985
- [Gopalakrishnan u. a. 2018] GOPALAKRISHNAN, Bharath ; SINGH, Arun K. ; KRISHNA, K. M. ; MANOCHA, Dinesh: Solving chance constrained optimization under nonparametric uncertainty through Hilbert space embedding. In: *arXiv preprint* (2018). – arXiv:1811.09311
- [Hodge u. Milligan 2011] HODGE, Bri-Mathias ; MILLIGAN, Michael: Wind power forecasting error distributions over multiple timescales. In: *Proceedings of the IEEE Power and Energy Society General Meeting*. Detroit, MI, 2011

- [Keil u. a. 2020] KEIL, Rachel E. ; MILLER, Alexander T. ; KUMAR, Mrinal ; RAO, Anil V.: *Method for chance constrained optimal control using biased kernel density estimators*. arXiv preprint. – arXiv:2003.08010
- [Keil u. a. 2021] KEIL, Rachel E. ; MILLER, Alexander T. ; KUMAR, Mrinal ; RAO, Anil V.: Method for solving chance constrained optimal control problems using biased kernel density estimators. In: *Optimal Control Applications and Methods* 42 (2021), Nr. 1, S. 330–354. <http://dx.doi.org/10.1002/oca.2675>. – DOI 10.1002/oca.2675
- [Keil u. Rao 2021] KEIL, Rachel E. ; RAO, Anil V.: Warm start method for solving chance constrained optimal control problems using biased kernel density estimators. In: *Journal of Dynamic Systems, Measurement, and Control* 143 (2021), Nr. 12, S. 124502. <http://dx.doi.org/10.1115/1.4051748>. – DOI 10.1115/1.4051748
- [Liu u.a. 2022] LIU, Wei ; YANG, Li ; YU, Bo: Kernel density estimation based distributionally robust mean-CVaR portfolio optimization. In: *Journal of Global Optimization* 84 (2022), S. 1053–1077. <http://dx.doi.org/10.1007/s10898-022-01177-5>. – DOI 10.1007/s10898-022-01177-5
- [Luedtke u. Ahmed 2008] LUEDTKE, James ; AHMED, Shabbir: A sample approximation approach for optimization with probabilistic constraints. In: *SIAM Journal on Optimization* 19 (2008), Nr. 2, S. 674–699. <http://dx.doi.org/10.1137/070702928>. – DOI 10.1137/070702928
- [Nemirovski u. Shapiro 2006] NEMIROVSKI, Arkadi ; SHAPIRO, Alexander: Convex approximations of chance constrained programs. In: *SIAM Journal on Optimization* 17 (2006), Nr. 4, S. 969–996. <http://dx.doi.org/10.1137/050622328>. – DOI 10.1137/050622328
- [Pagnoncelli u. a. 2009] PAGNONCELLI, Bernardo K. ; AHMED, Shabbir ; SHAPIRO, Alexander: Sample average approximation method for chance constrained programming: theory and applications. In: *Journal of Optimization Theory and Applications* 142 (2009), Nr. 2, S. 399–416. <http://dx.doi.org/10.1007/s10957-009-9523-6>. – DOI 10.1007/s10957-009-9523-6
- [Parzen 1962] PARZEN, Emanuel: On estimation of a probability density function and mode. In: *Annals of Mathematical Statistics* 33 (1962), Nr. 3, S. 1065–1076. <http://dx.doi.org/10.1214/aoms/1177704472>. – DOI 10.1214/aoms/1177704472

- [Peña-Ordieres u. a. 2020] PEÑA-ORDIERES, Alejandra ; LUEDTKE, James R. ; WÄCHTER, Andreas: Solving chance-constrained problems via a smooth sample-based nonlinear approximation. In: *SIAM Journal on Optimization* 30 (2020), Nr. 3, S. 2221–2250. <http://dx.doi.org/10.1137/19M1261985>. – DOI 10.1137/19M1261985
- [Prékopa 1995] PRÉKOPA, András: *Stochastic Programming*. Dordrecht : Kluwer Academic Publishers, 1995
- [Qi u. a. 2026] QI, Xiaotong ; LIU, Wei ; YU, Bo ; YANG, Li: *Distributionally robust chance-constrained optimal load shedding model for active distribution networks based on KDE*. Optimization Online preprint, 2026. – Optimization Online ID 34324
- [Rockafellar u. Uryasev 2000] ROCKAFELLAR, R. T. ; URYASEV, Stanislav: Optimization of conditional value-at-risk. In: *Journal of Risk* 2 (2000), Nr. 3, S. 21–41. <http://dx.doi.org/10.21314/JOR.2000.038>. – DOI 10.21314/JOR.2000.038
- [Rosenblatt 1956] ROSENBLATT, Murray: Remarks on some nonparametric estimates of a density function. In: *Annals of Mathematical Statistics* 27 (1956), Nr. 3, S. 832–837. <http://dx.doi.org/10.1214/aoms/1177728190>. – DOI 10.1214/aoms/1177728190
- [Savage 2009] SAVAGE, Sam L.: *The Flaw of Averages: Why We Underestimate Risk in the Face of Uncertainty*. Hoboken, NJ : Wiley, 2009
- [Schuster 2025] SCHUSTER, Michael: On the convergence of optimization problems with kernel density estimated probabilistic constraints. In: *Applied Set-Valued Analysis and Optimization* 7 (2025), Nr. 2, S. 209–221
- [Shapiro u. a. 2014] SHAPIRO, Alexander ; DENTCHEVA, Darinka ; RUSZCZYŃSKI, Andrzej: *Lectures on Stochastic Programming: Modeling and Theory*. 2. Philadelphia, PA : SIAM, 2014. <http://dx.doi.org/10.1137/1.9781611973433>. <http://dx.doi.org/10.1137/1.9781611973433>
- [Silverman 1986] SILVERMAN, Bernard W.: *Density Estimation for Statistics and Data Analysis*. London : Chapman & Hall, 1986
- [Soyster 1973] SOYSTER, Allen L.: Convex programming with set-inclusive constraints and applications to inexact linear programming. In: *Operations Research* 21 (1973), Nr. 5, S. 1154–1157. <http://dx.doi.org/10.1287/opre.21.5.1154>. – DOI 10.1287/opre.21.5.1154

- [Wan u. a. 2020] WAN, Can ; ZHAO, Changfei ; SONG, Yonghua: Chance constrained extreme learning machine for nonparametric prediction intervals of wind power generation. In: *IEEE Transactions on Power Systems* 35 (2020), Nr. 5, S. 3869–3884. <http://dx.doi.org/10.1109/TPWRS.2020.2986282>. – DOI 10.1109/TPWRS.2020.2986282
- [Wang u. a. 2021] WANG, Jingyao ; WANG, Cheng ; LIANG, Yile ; BI, Tao ; SHAFIE-KHAH, Miadreza ; CATALÃO, João P. S.: Data-driven chance-constrained optimal gas-power flow calculation: a Bayesian nonparametric approach. In: *IEEE Transactions on Power Systems* 36 (2021), Nr. 5, S. 4683–4698. <http://dx.doi.org/10.1109/TPWRS.2021.3065465>. – DOI 10.1109/TPWRS.2021.3065465
- [Wu u. Kargarian 2023] WU, Chutian ; KARGARIAN, Amin: Computationally efficient data-driven joint chance constraints for power systems scheduling. In: *IEEE Transactions on Power Systems* 38 (2023), Nr. 3, S. 2858–2869. <http://dx.doi.org/10.1109/TPWRS.2022.3195127>. – DOI 10.1109/TPWRS.2022.3195127
- [Wu u. a. 2021] WU, Chutian ; KARGARIAN, Amin ; JEON, Hyun W.: Data-driven nonparametric joint chance constraints for economic dispatch with renewable generation. In: *IEEE Transactions on Industry Applications* 57 (2021), Nr. 6, S. 6537–6546. <http://dx.doi.org/10.1109/TIA.2021.3105364>. – DOI 10.1109/TIA.2021.3105364

Erklärung

Hiermit versichere ich, dass ich die vorliegende Arbeit selbstständig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel benutzt habe, dass alle Stellen der Arbeit, die wörtlich oder sinngemäß aus anderen Quellen übernommen wurden, als solche kenntlich gemacht sind und dass die Arbeit in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegt wurde.

Erlangen, den June 14, 2026

Hierher die Unterschrift

LEBENS LAUF

VORNAME NACHNAME

ZUR PERSON

Name	Vorname Nachname
Geboren	
Adresse	Straße Hausnummer
	PLZ Nürnberg

SCHULE

09/1977 - 08/1981	1. Schulzeit
09/1981 - 06/1990	2. Schulzeit

STUDIUM

09/1991 - 08/1993	1. Universit"at
09/1993 - 06/1996	2. Universit"at

SONSTIGE T"ATIGKEITEN

09/1990 - 08/1991	Anlauf und Sprung
-------------------	-------------------
